



# Linux Threads

Kiran Divekar

Nevis Networks, Pune

[kirandivekar@gmail.com](mailto:kirandivekar@gmail.com)

[www.geocities.com/kirandivekar/](http://www.geocities.com/kirandivekar/)



**G E E P**  

---

**GEEKS OF PUNE**



# Abstract

- What are threads? User space
- Threads Basics
- Creation and termination: `thread_info`
- User space thread development
- Threads Debugging
- Thread Experience
- Kernel Threads
- Latest Kernel Developments



# User space threads

- Threads are separate contexts of execution within the same process.
- Fork(), vfork(), clone() create a process
- Threads share the same global memory space, but they have different stacks.
- `<linux/sched.h>` task\_struct



# User space threads

- Process creation: `clone()`.
- Kernel : `do_fork`
- Threads creation: `pthread_create`.
- Kernel : `do_fork`



# Thread\_info

```
Struct thread_info {  
    struct task_struct *task ;  
    unsigned long flags ;  
  
    exec_domain ;  
    cpu ;  
    preempt_count ;  
    add_limit ;  
    supervisor_stack ;  
}
```



# Copy On Write [COW]

- Linux threads follow COW technique
- Process descriptor is duplicated if written.
- CLONE\_VM
- CLONE\_FILES
- CLONE\_FS
- CLONE\_SIGHAND
- Note: Child execs first. Why ?
- Shared page tables : kernel 2.7 ???





# Threads Exit

- User space : `exit` / `pthread_cancel`
- Kernel : `do_exit()`.
- Recycling pids as per MAX value
- Remember : `thread_info`
- Threads are scheduled.
- Why 2 separate structures ?



# Thread Development

- Pthreads
  - pthread\_create, pthread\_attributes.
  - Create a detached thread
- Basic Linuxthreads
  - Why thread-safe, reentrant ?
  - Example: errno, strtok
- TLS
- NPTL





# Thread Local Storage

- The purpose of TLS is to give each thread access to a region of memory which is *not* shared with all other threads. It is stored just below the stack address.
  - `__thread int i;`
  - `static __thread char *p;`
- When the `addr-of` operator is applied to a thread-local variable, it returns the address of the current thread's instance of that variable at run time.
- `/lib/tls/libc.so`
- Example : can;t find TLS link



# Disadvantages

- The manager thread control: performance
- Issues with SMP systems
- Requirement of pthread\_exit
- Context switching delay
- N+2 threads
- POSIX incompatible
  
- NPTL



# Native POSIX Thread Library

- No manager thread
- Kernel scheduled
- POSIX synchronization using futex
- Per process signal handling rather than per thread
- ABI support
- Backward compatibility using `LD_ASSUME_KERNEL`
- `getconf GNU_LIBPTHREAD_VERSION`



# Kernel Requirements

- Ingo Molnar: Early 2003
- Support for an arbitrary number of thread-specific data areas
- Removes threads per process limit
- Clone system call extended to optimize creation and facilitate termination of threads [ detached thread ]
- Futex wakeup on thread Id [ pthread\_join ]
- Signals sent per thread. Fatal signals terminate process
- /proc contains only process information



# Thread Debugging

## Fork Behaviour

- Default :
  - `set follow-fork-mode parent`
- Child breakpoint gets a SIGTRAP
- Separate gdb to debug child

## ➤ Pthread behaviour

- `gdb 6 & above`
- `info threads`
- `thread apply [all / threadno] cmd`





# Thread experience

- Separate open system call
- Crash analysis.
  - Kernel support for thread backtrace
  - Critical section analysis
- Thread scheduling
  - Unlocking thread wakes up one thread
- Blocking behaviour
  - e.g. system call in an embedded system
- Avoid asynchronous thread cancellation
  - It may be in critical section





# Thread experience

- Timing issues
  - Socket descriptor manipulation
- Any more ?



# Kernel execution

- Kernel Context
  - Asynchronous: Interrupt handlers
- Process [User] Context
  - Synchronous: Response to system calls
- Kernel threads
  - Similar to user space daemons.



# Kernel threads

ps -ef

ID	PID	PPID	STIME	TIME	CMD
----	-----	------	-------	------	-----

➤ root	1	0	22:36	00:00:00	init [3]
➤ root	2	1	22:36	00:00	[ksoftirqd/0]
root	3	1	22:36	00:00:00	[events/0]
root	38	3	22:36	00:00:00	[pdflush]
root	39	3	22:36	00:00:00	[pdflush]
root	29	1	22:36	00:00:00	[khubd]
root	695	1	22:36	00:00	[kjournald]
root	3914	1	22:37	00:00:00	[nfsd]
root	3915	1	22:37	00:00:00	[nfsd]
root	4066	4015	22:59	00:00:00	ps -ef



# Kernel Threads

- The [ksoftirqd/0] kernel thread is an aid to implement soft IRQs
- The events/n threads ( $n = \text{processor number}$ ) help implement work queues
- The pdflush kernel thread flushes dirty pages from the page cache
- The khubd thread, part of the Linux USB core, monitors the machine's USB hub
- The nfsd thread, monitoring network filesystems



# Why Kernel Threads?

Candidates for a kernel thread because:

- It's a background task, since it has to wait for asynchronous events.
  - It needs access to kernel data structures, since the actual detection of events must be done by other parts of the kernel.
  - It has to invoke a user-mode helper program, which is a time consuming operation.
- 
- Ex: monitor receive buffers of a network
  - `create_kthread` ???





# Kernel Thread How?

- Kernel thread Life Cycle:
  - Daemonize
  - Make init as my parent
  - Wait for event using wait queue.
- Event Generation
  - Kernel data structure monitoring
  - If the health is unsatisfactory, ( eg. Buffers below low watermark), wake up the queue.
- When event comes, check for signal SIGKILL
  - Start kthread operation
  - Set state to TASK\_RUNNING





# Kernel threads...

- User space interaction

**SYSCTL** : system control operations.

- `/proc/sys/kernel/modprobe`
- `/proc/sys/kernel/hotplug`

**Reference:**

`ksoftirqd` : `kernel/softirq.c`

`pdflush` : `mm/pdflush.c`

`khubd` drivers : `usb/core/hub.c`



# Latest Kernel Development

- Bottom Halves: Interrupt Deferred routine
- Tasklet. Deferred work processing
- Task queues: Fixed. TQ\_INTERRUPT
- Work queues: Event Handling
  
- Syslets
- Threadlets



# Syslets...

- Syslets are small, simple, lightweight programs (consisting of system-calls, 'atoms') that the kernel can execute autonomously (and, not the least, asynchronously), without having to exit back into user-space.
- Makes use of cachemiss technique.
- There are open issues.



# Threadlets...

- "Threadlets" are basically the user-space equivalent of syslets: small functions of execution that the kernel attempts to execute without scheduling. If the threadlet blocks, the kernel creates a real thread from it, and execution continues in that thread.
- The 'head' context (the context that never blocks) returns to the original function that called the threadlet.



# Happy Kernel Hacking !!!

## THANK YOU



**G E E P**  

---

**GEEKS OF PUNE**